

Red Hat OpenShift Container Platform


- Public/private DC.
- Bare metal and multiple cloud and virtualization providers.
- Full control by customer.

Red Hat OpenShift Dedicated

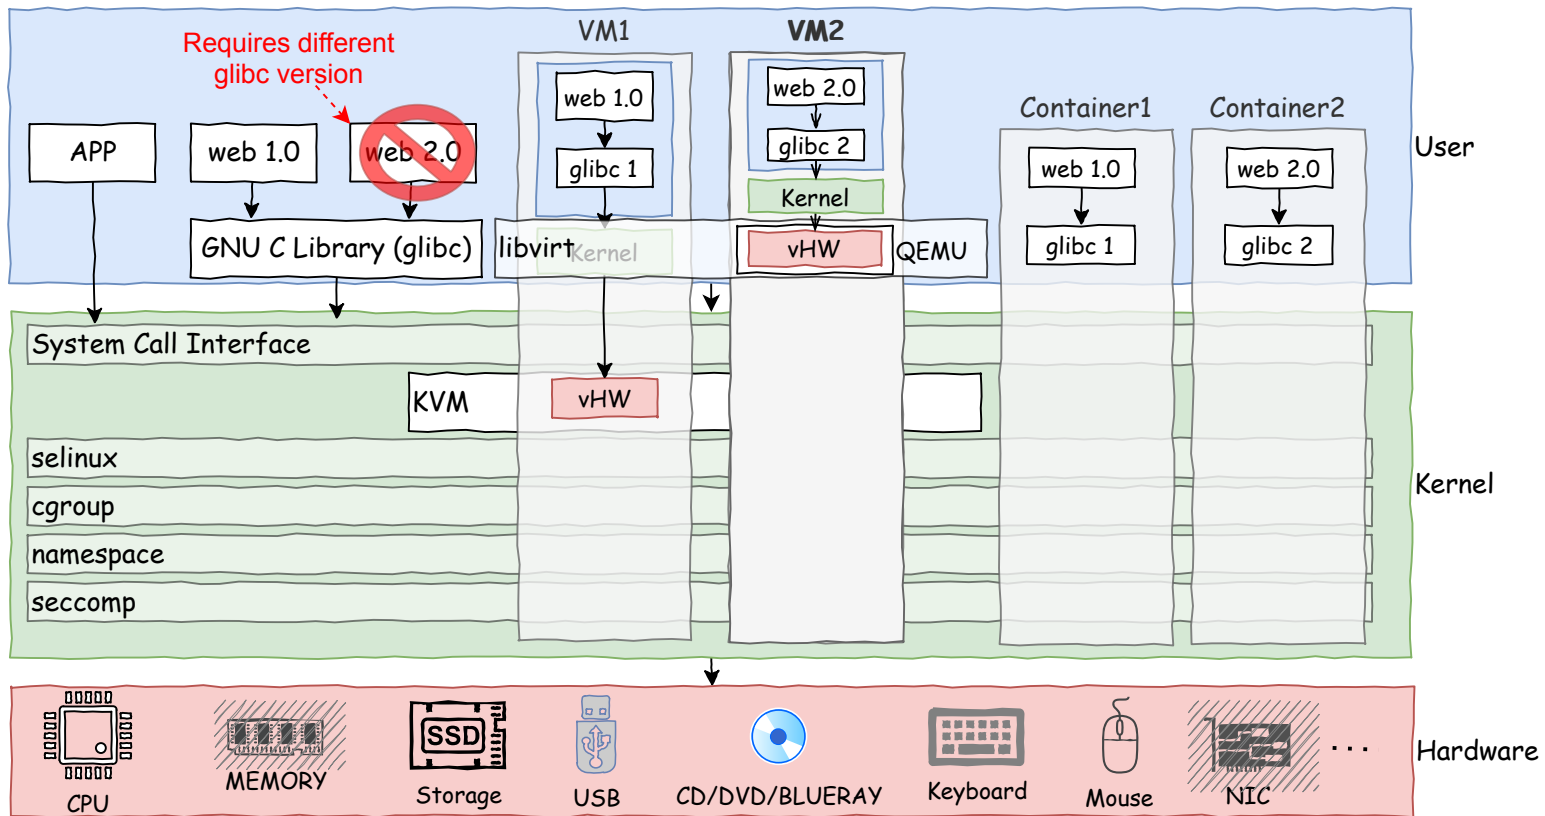
- Managed cluster in public cloud.
- RH manages the cluster.
- Customer manages updates and add-on services.

Red Hat OpenShift Online


- Public hosted cluster.
- Shared resources by multiple customers.
- RH manages cluster life cycle.



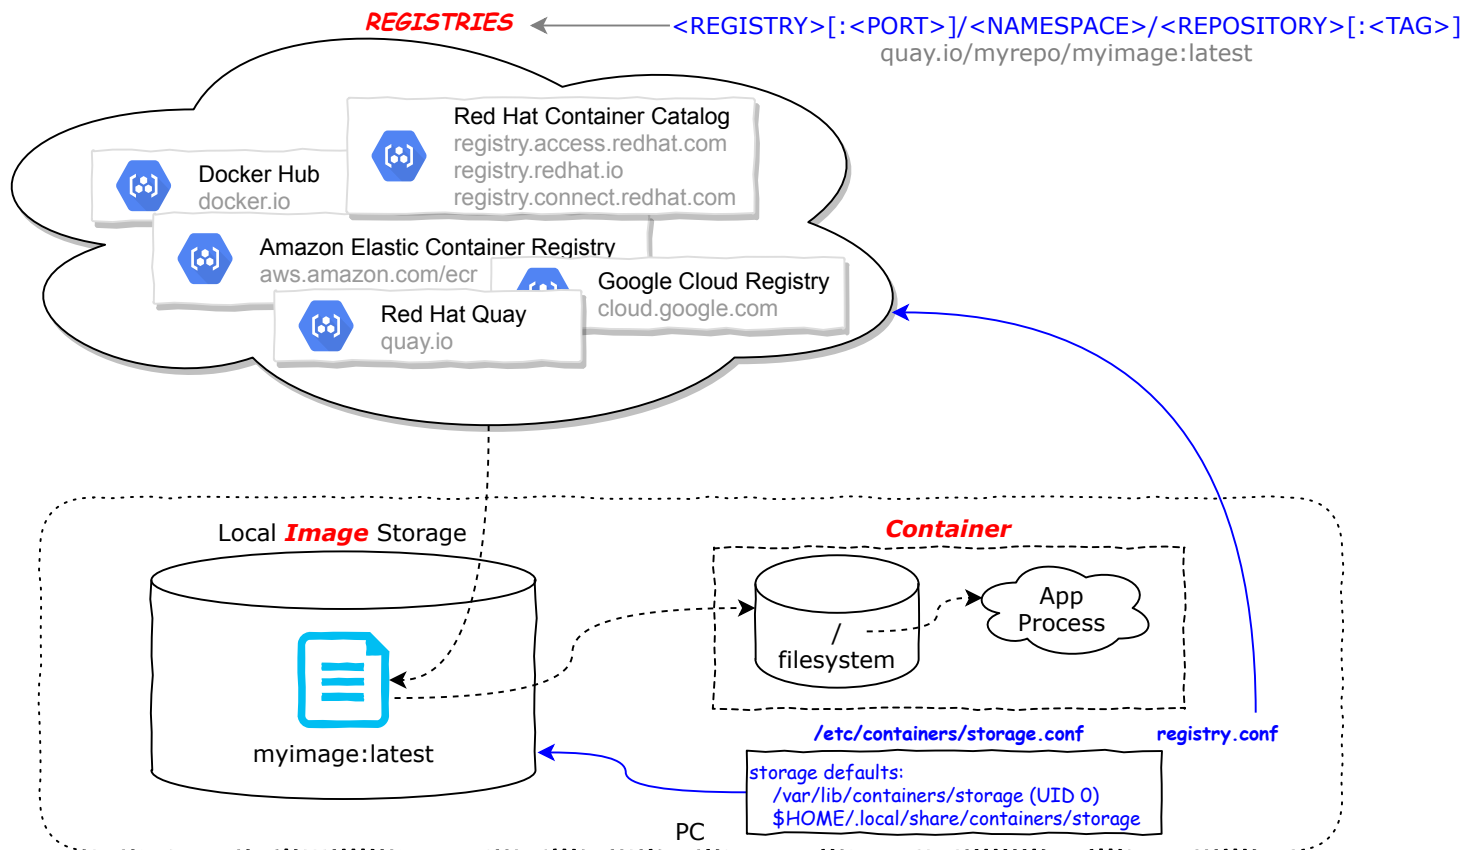
VM vs Container



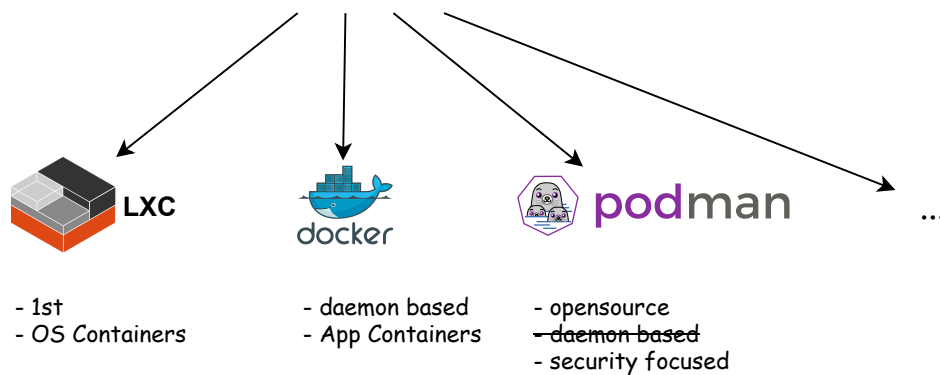
Ref: <https://www.redhat.com/en/blog/all-you-need-know-about-kvm-userspace>
<https://www.packetcoders.io/what-is-the-difference-between-qemu-and-kvm/>



Container Architecture



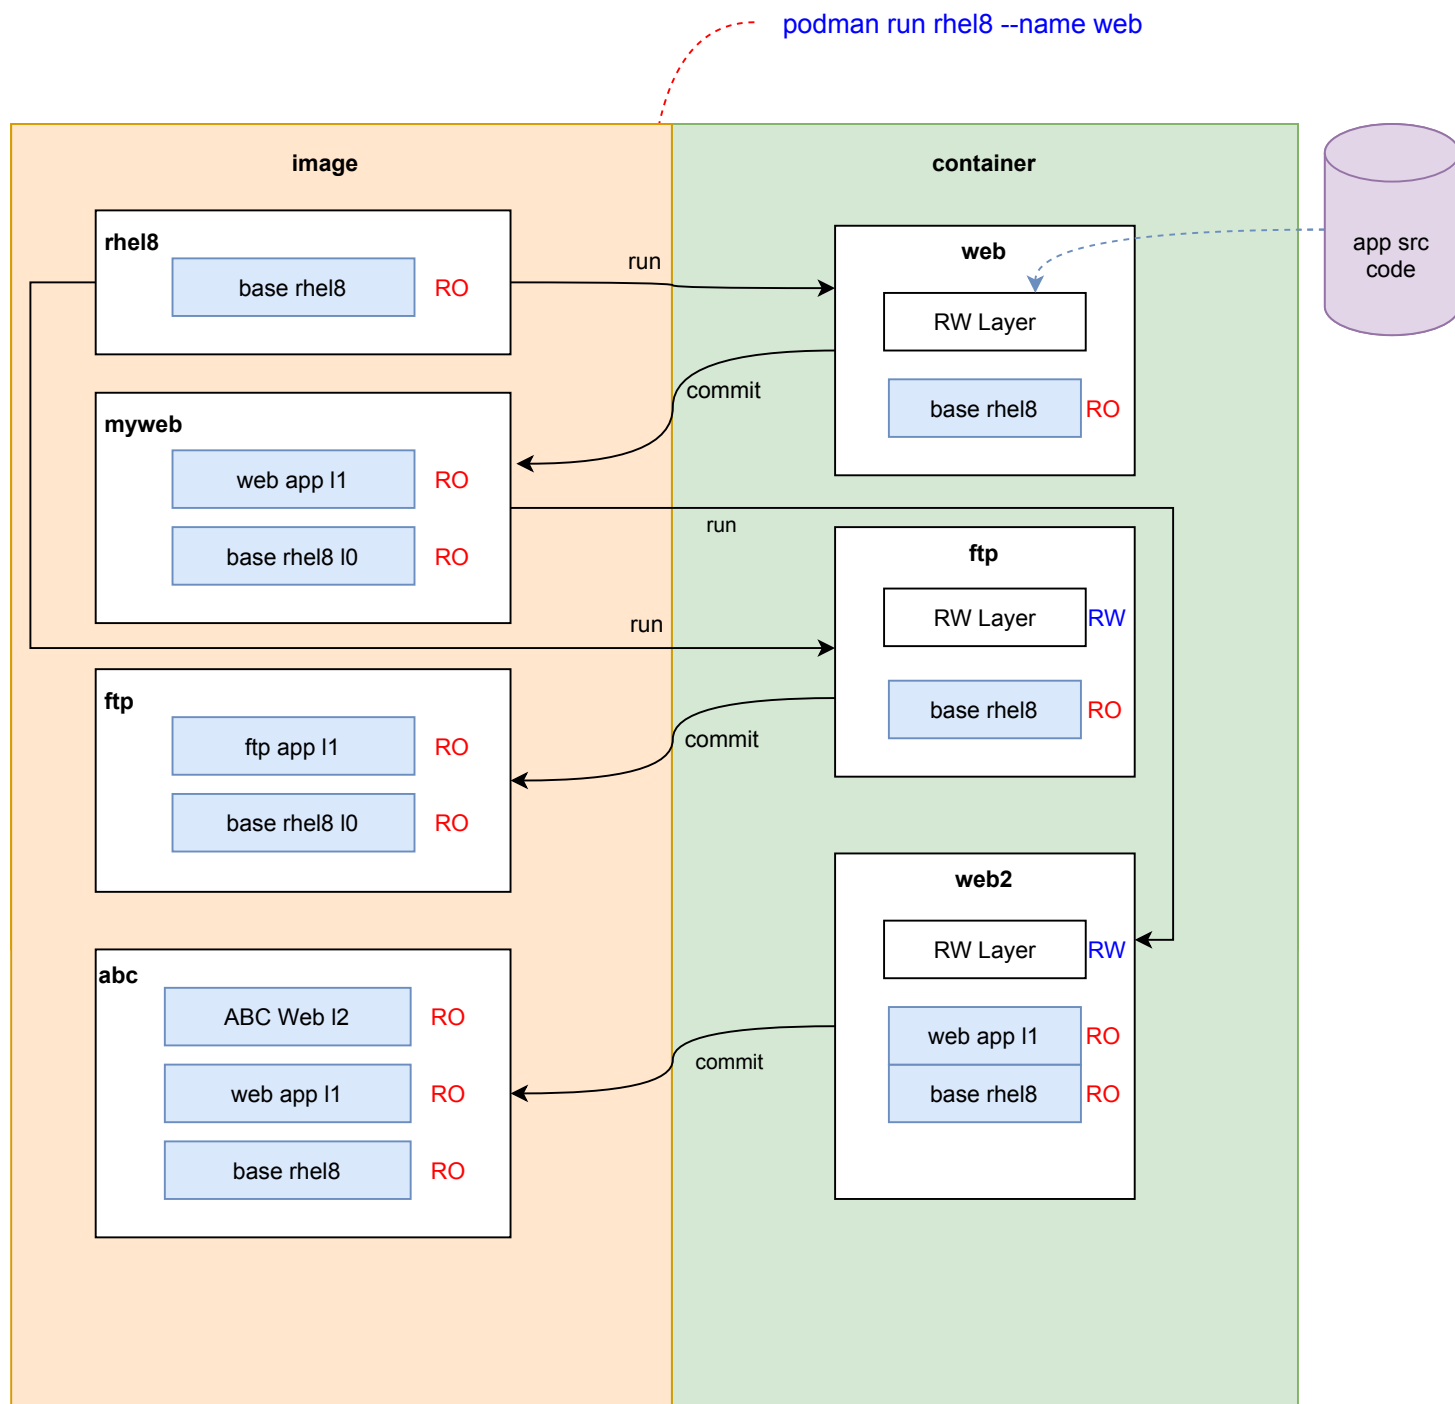
Container Utilities



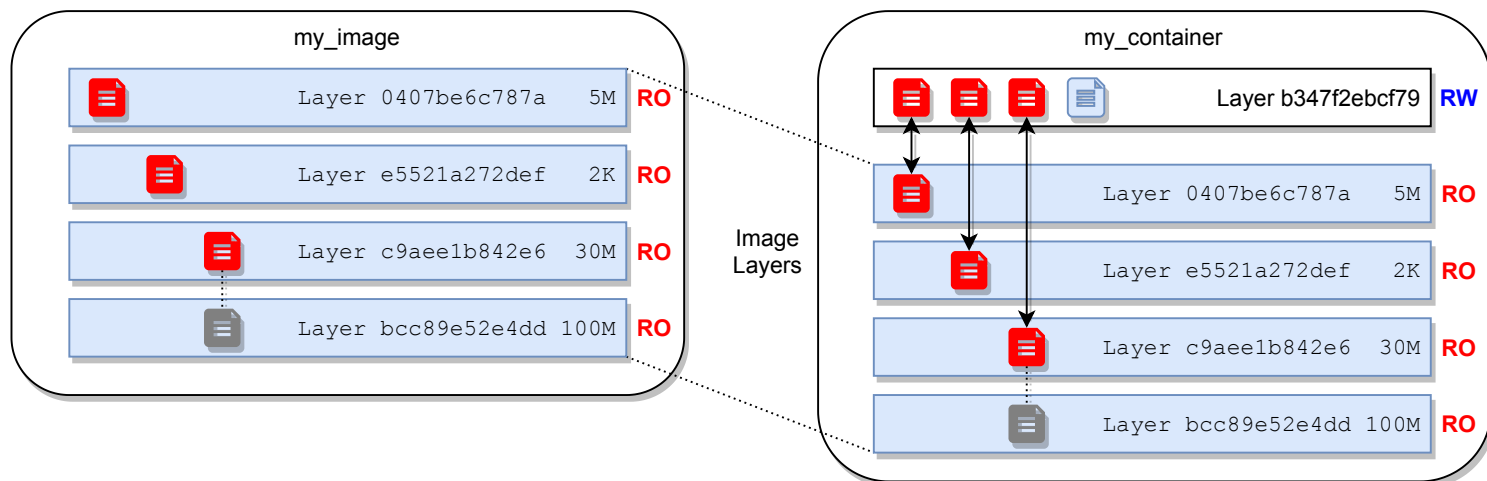
OS Container Vs Application Containers

Creating Image

1. Manual
2. Dockerfile/Containerfile
3. Source-To-Image(s2i/STI)
 - a) get runtime image and create container
 - b) clone source code into container
 - c) compile source code
 - d) deploy/publish compiled app
 - e) cleanup
 - f) save container as image

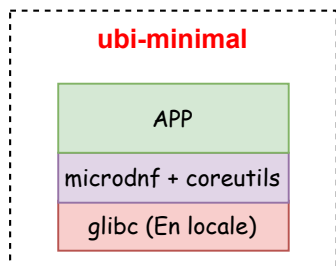


UnionFS - A Stackable Unification File System



BASE IMAGE TYPES


MINIMAL



Designed for apps that contain their own dependencies (Python, Node.js, .NET, etc.)

- Minimized pre-installed content set
- no suid binaries
- minimal pkg mgr (install, update & remove)


PLATFORM



For any apps that runs on RHEL

- Unified, OpenSSL crypto stack
- Full YUM stack
- Includes useful basic OS tools (tar, gzip, vi, etc)

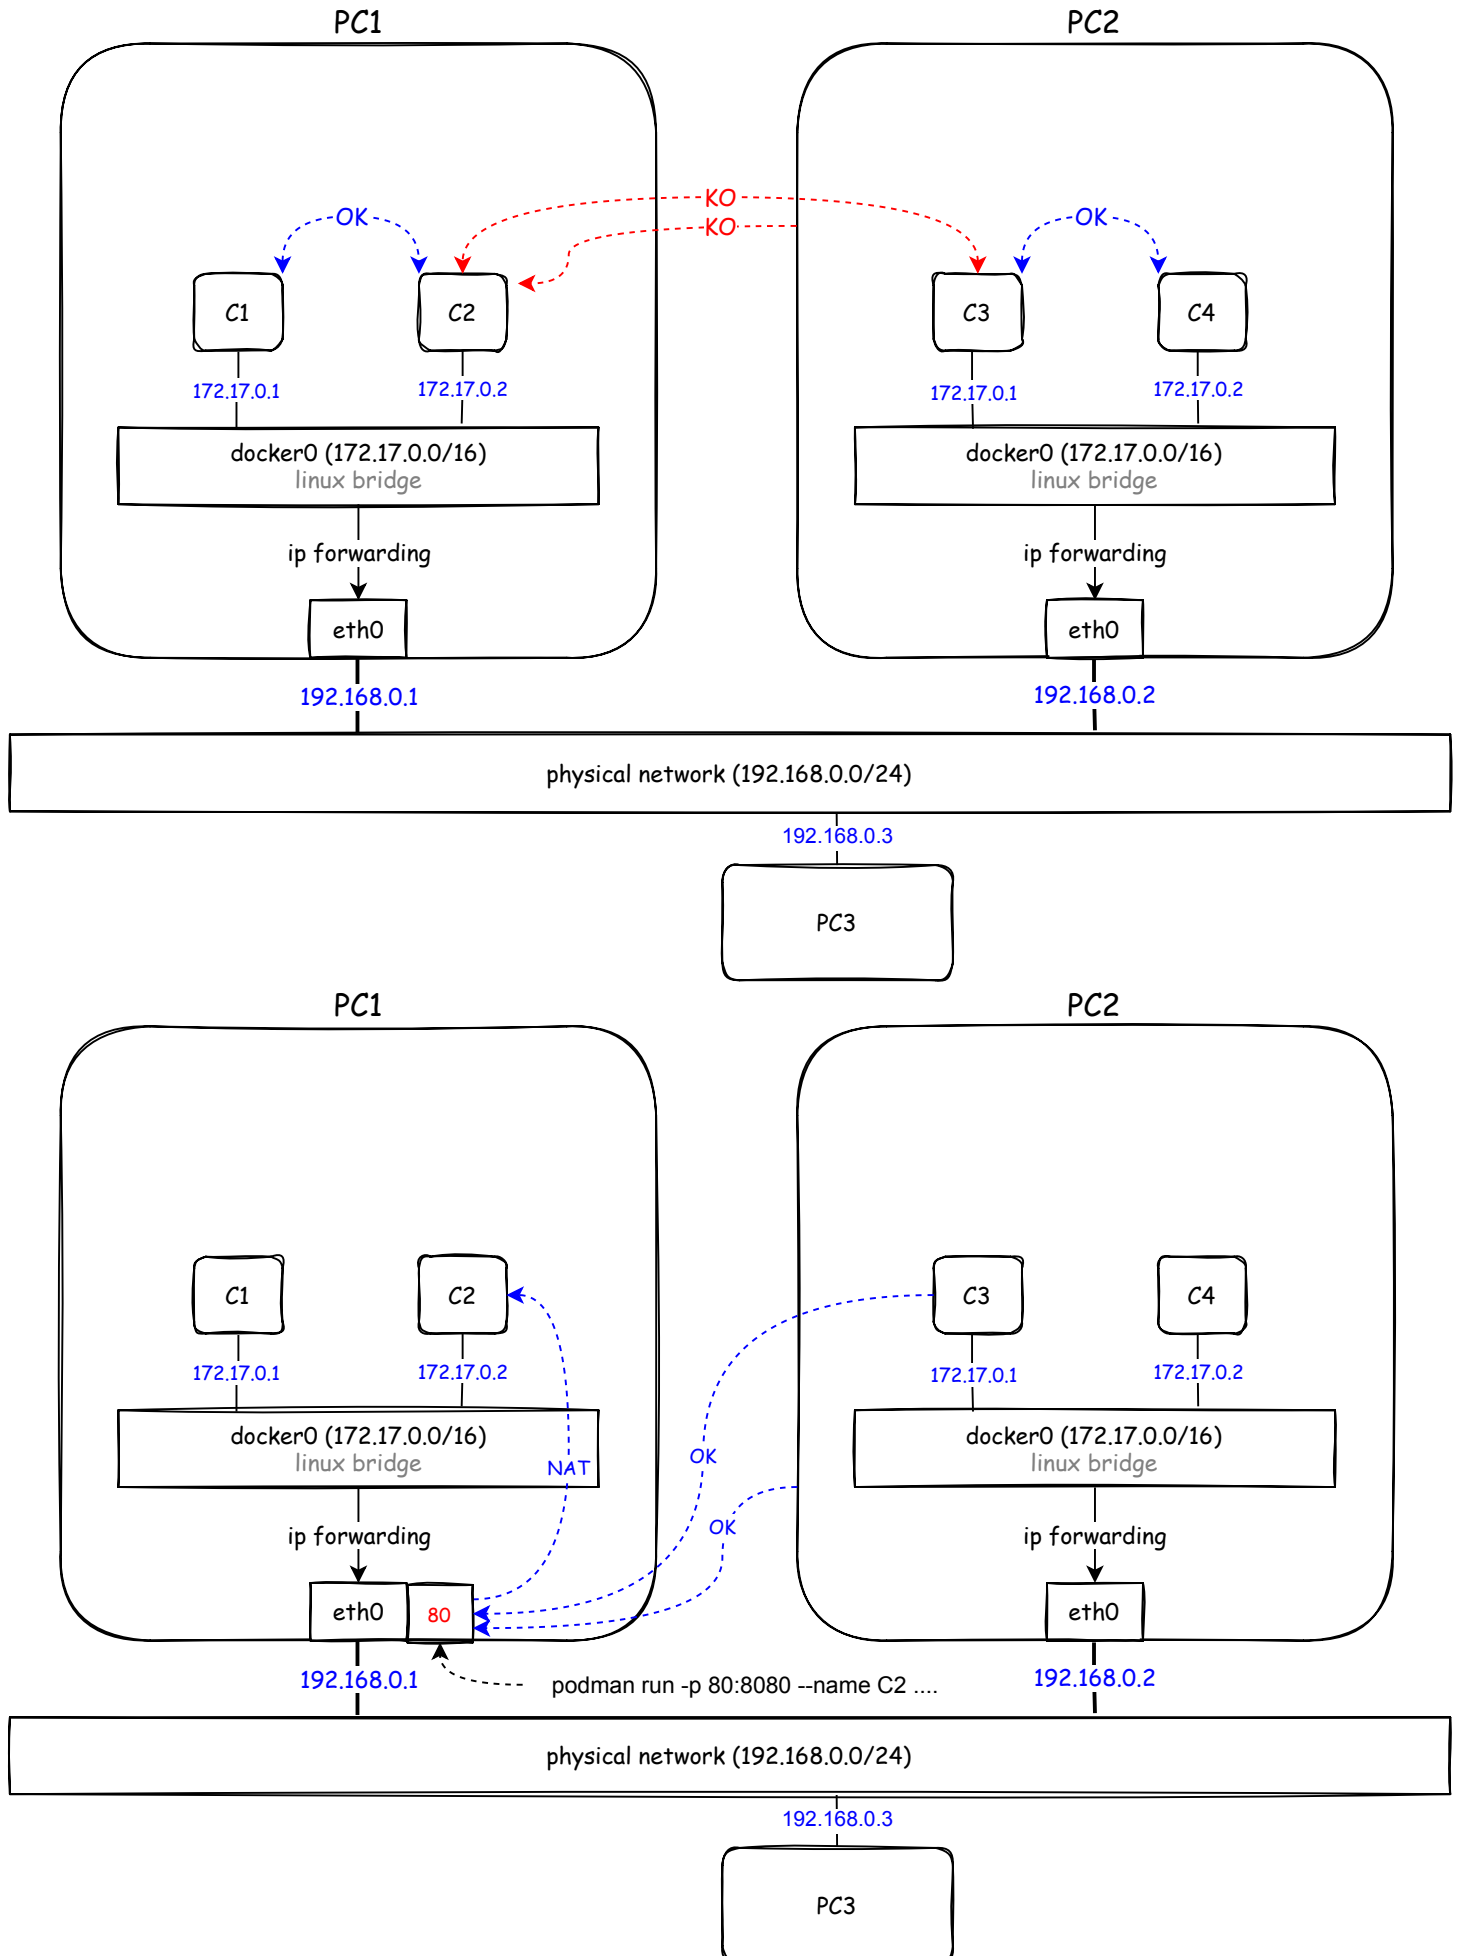
MULTI-SERVICE

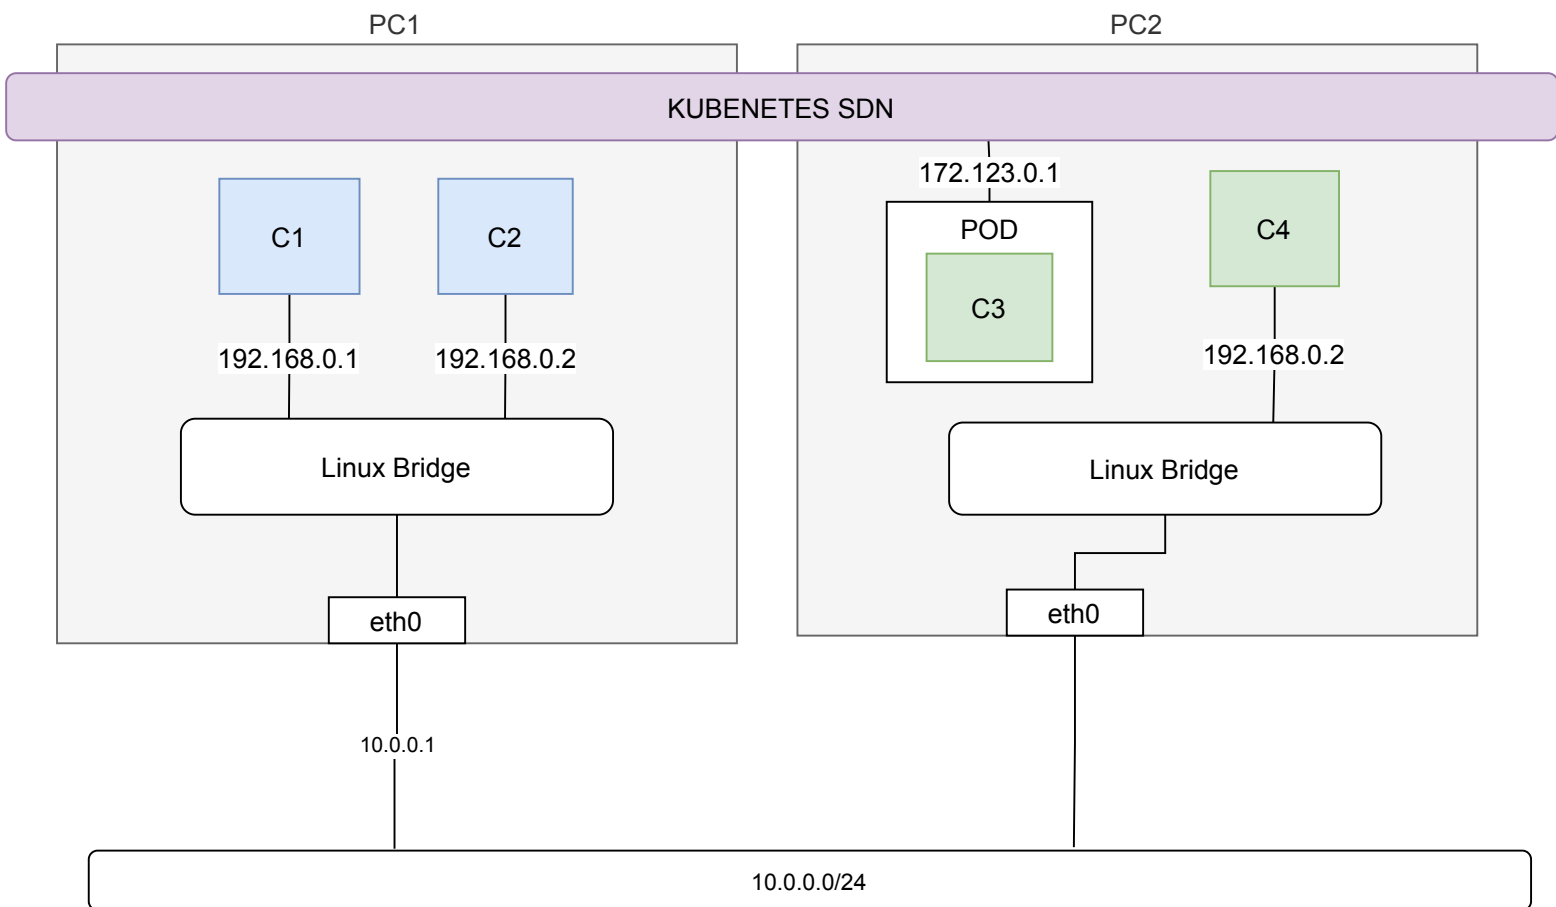
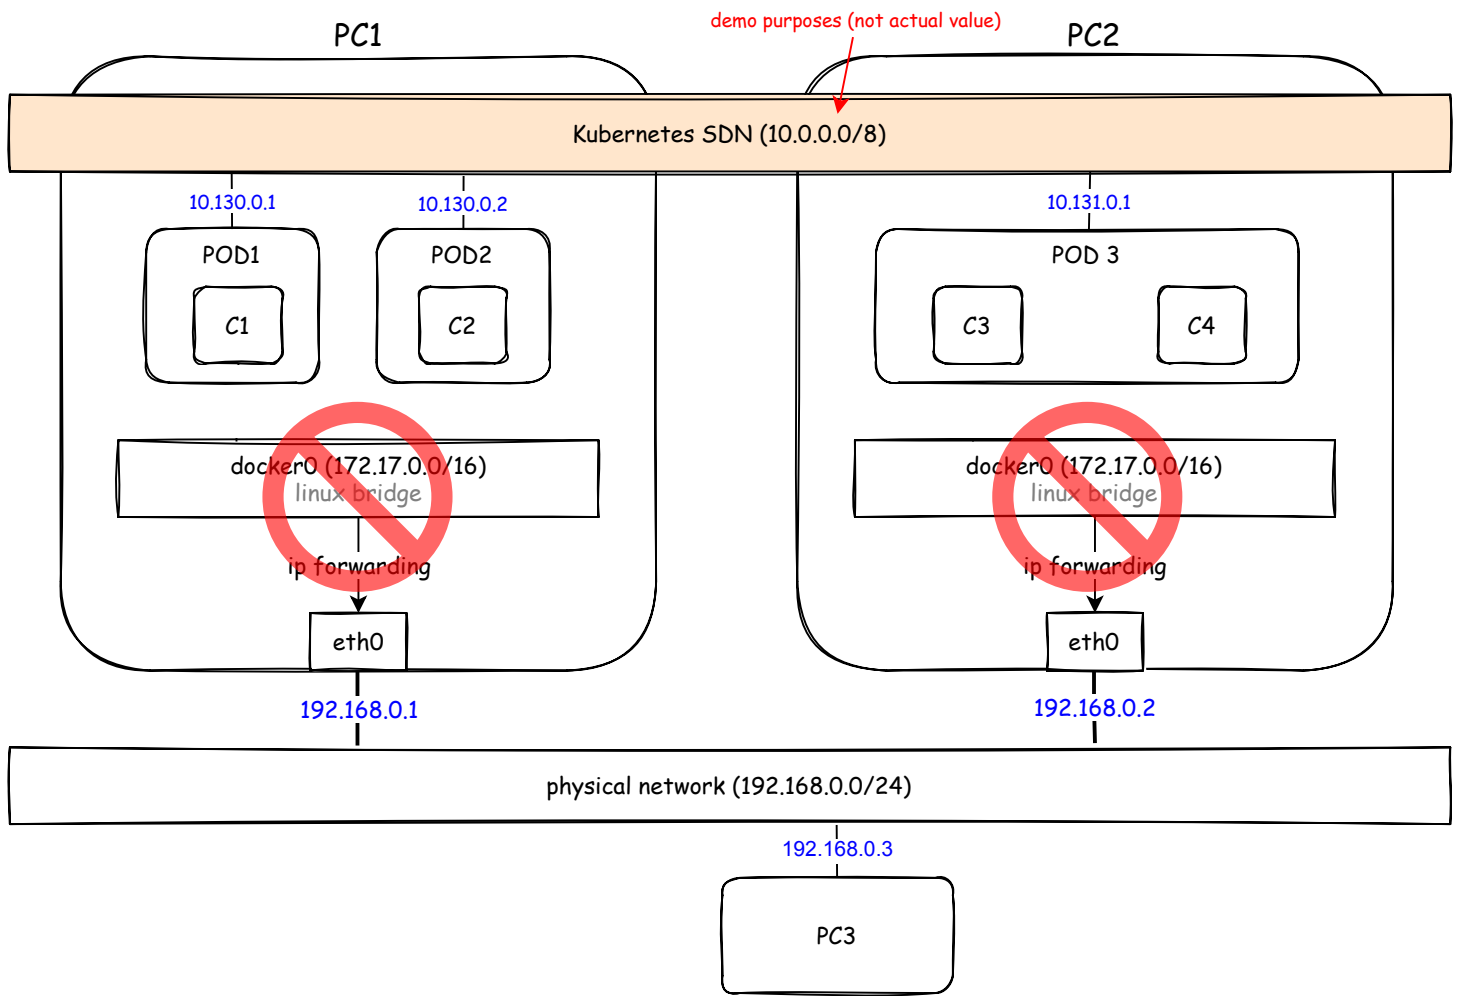


Eases running multi-service in single container

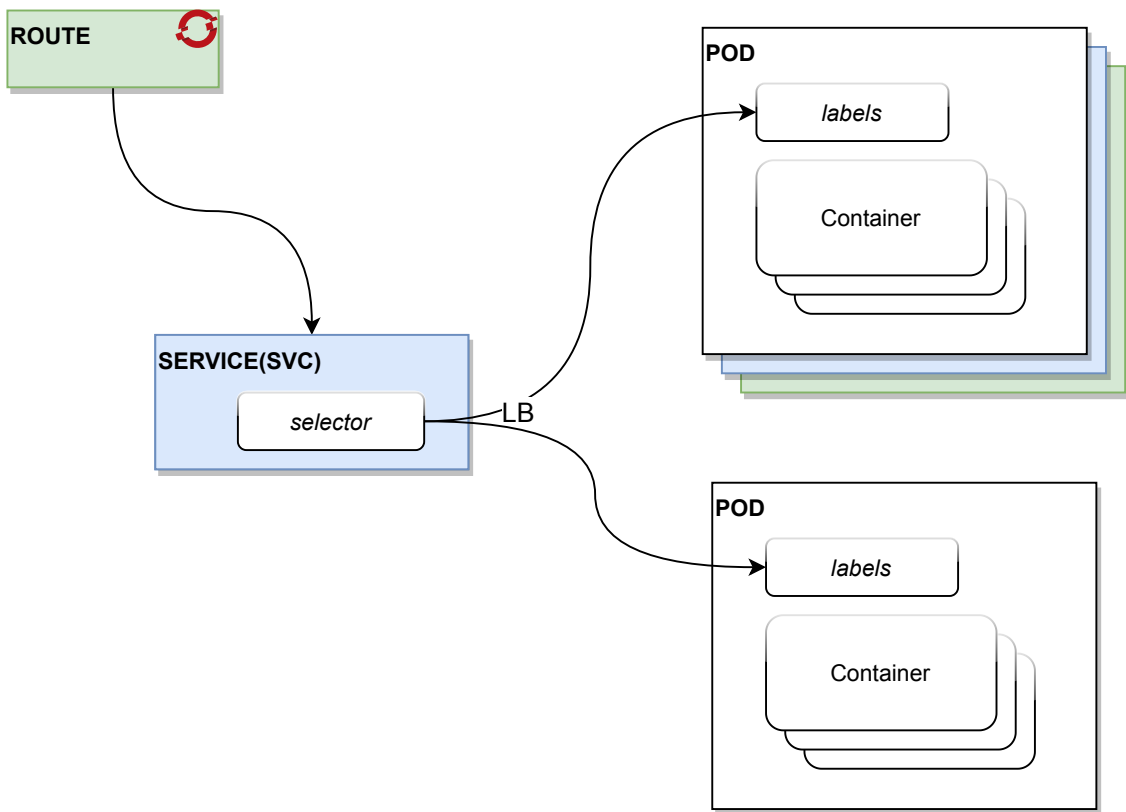
- configured to run systemd on start
- allows you to enable th services at build time

Basic Network - Container vs Kubernetes





Route, Service and Pod Relationship



POD

A pod contains one or more containers.

SERVICE

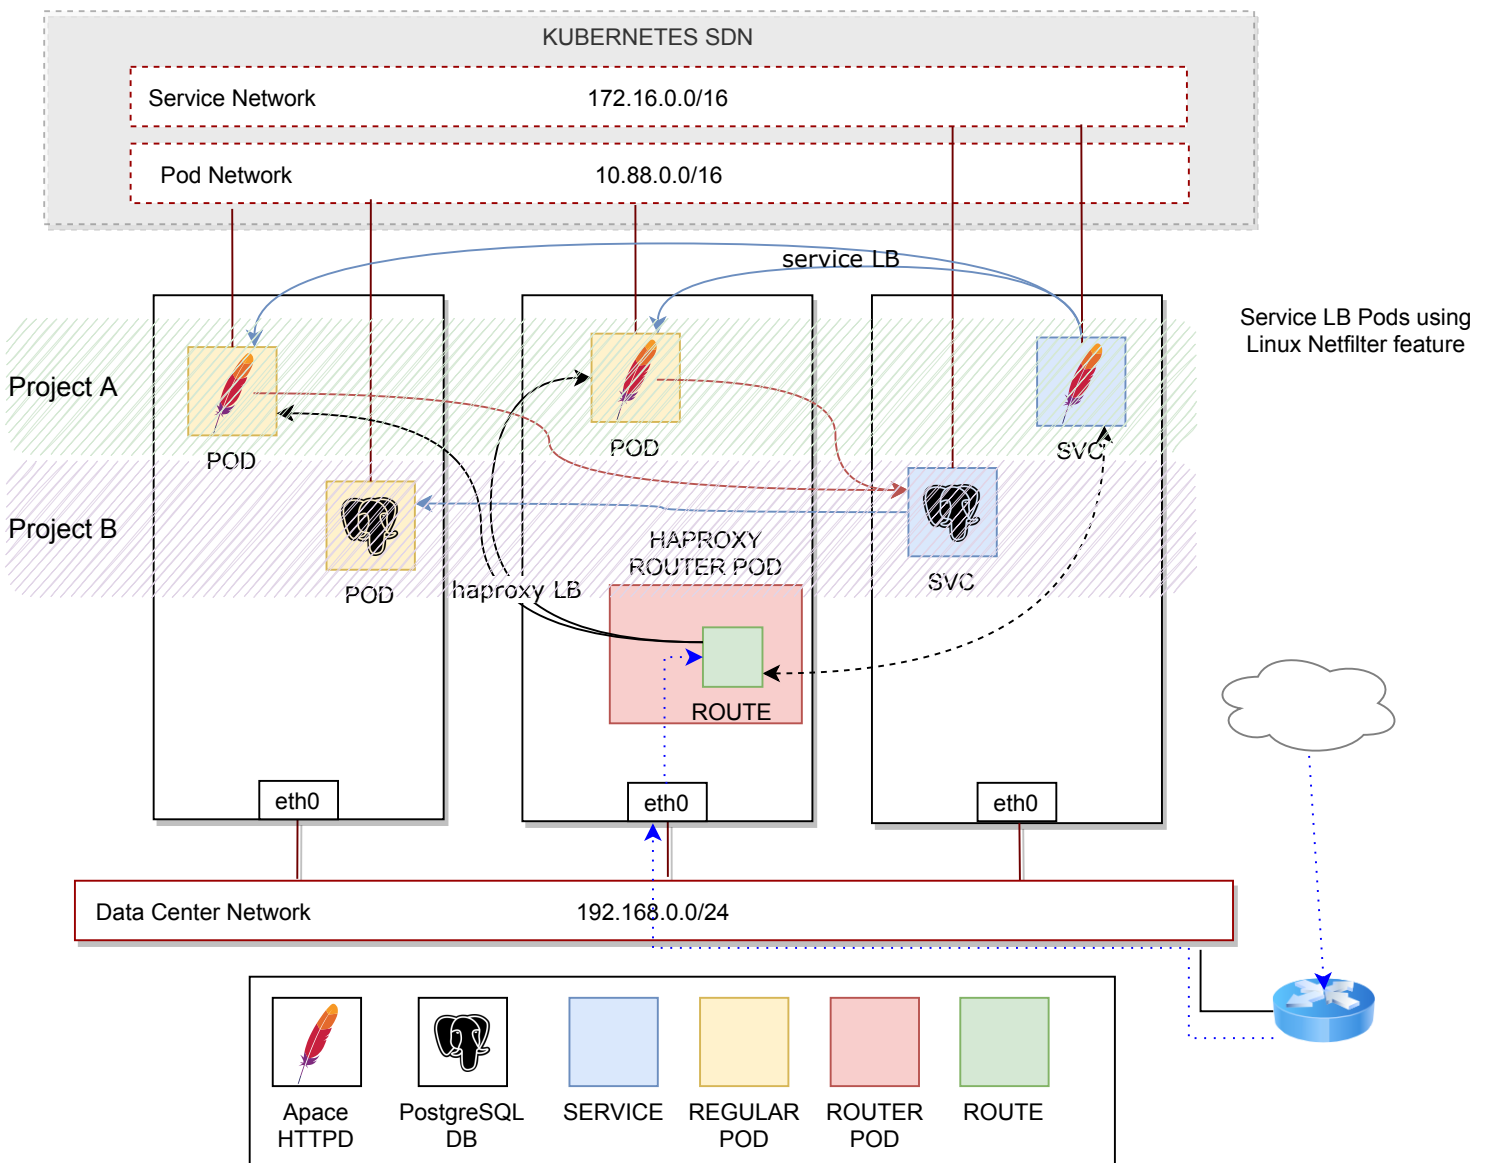
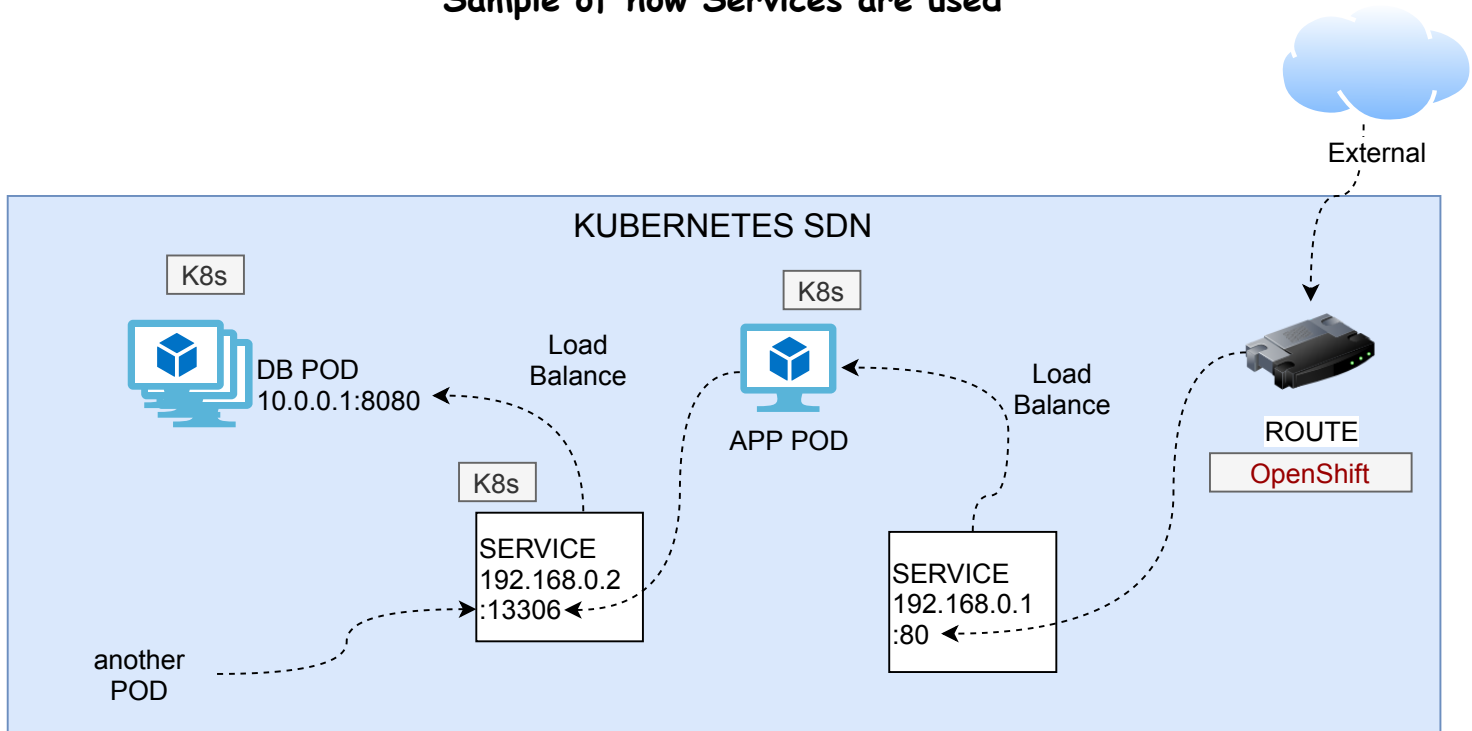
A service references the pod(s) by using the label selector.
The service load balances the connections between all the pods.

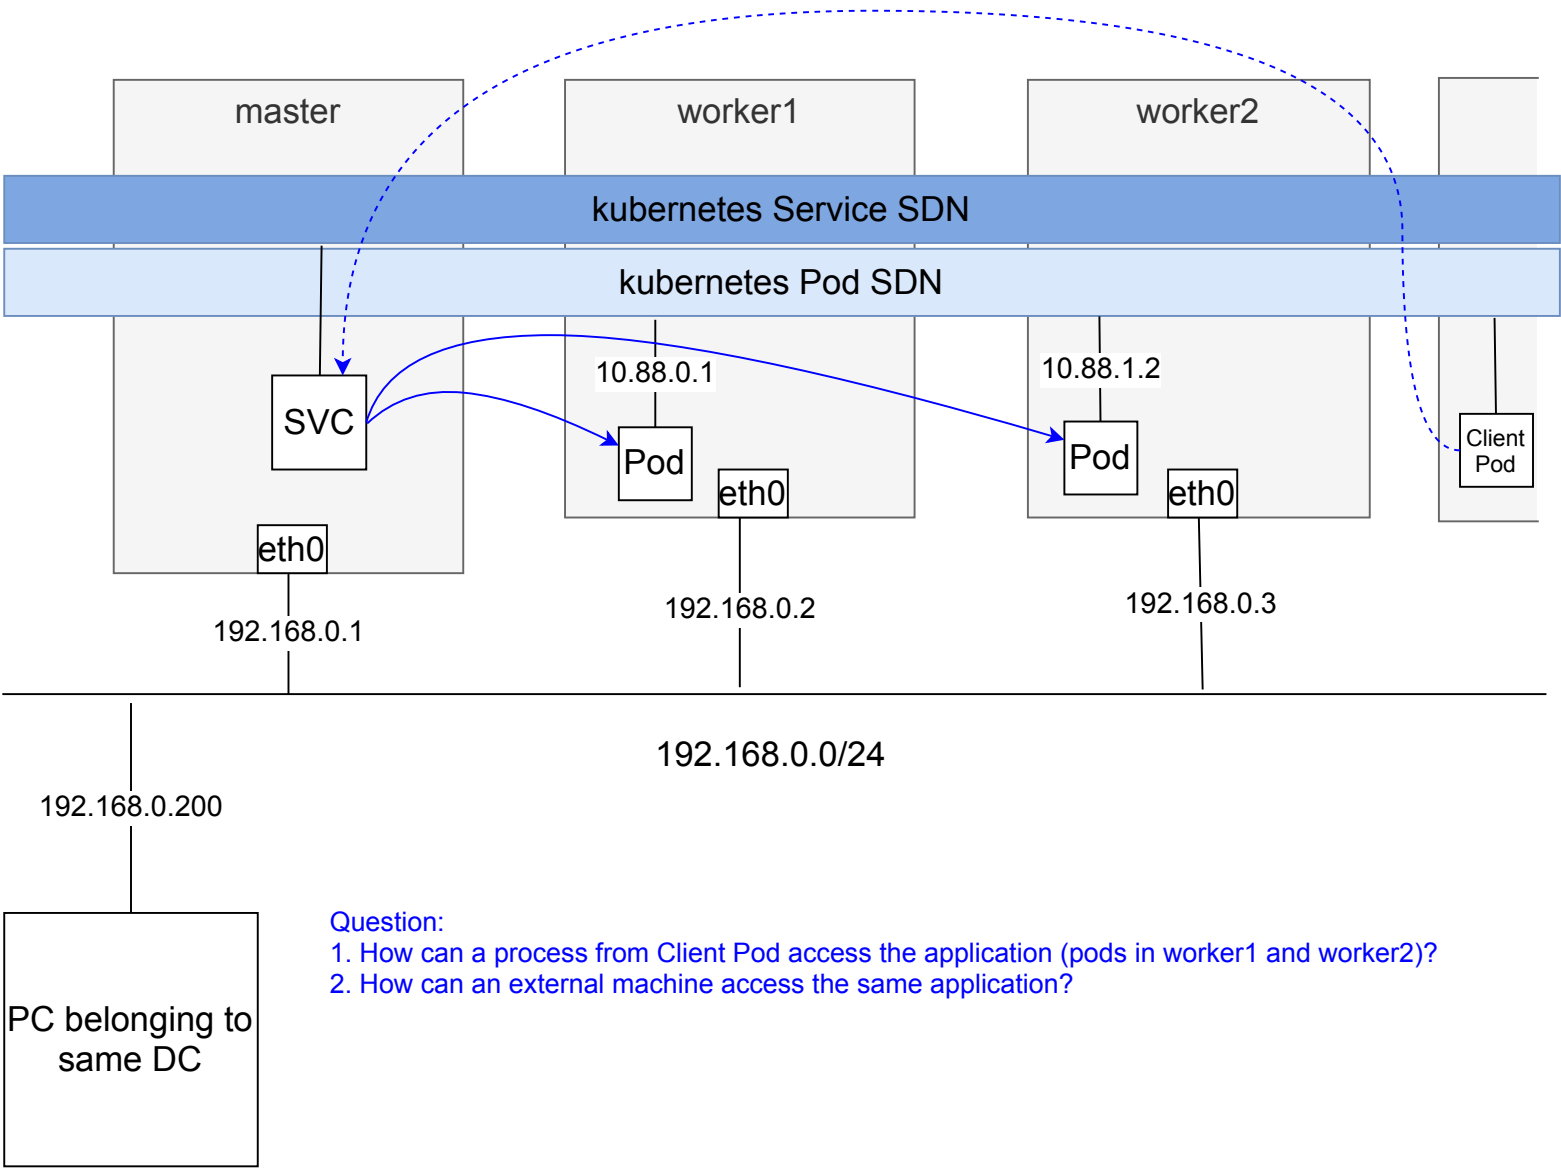
ROUTE

A route exposes the service to the external world.

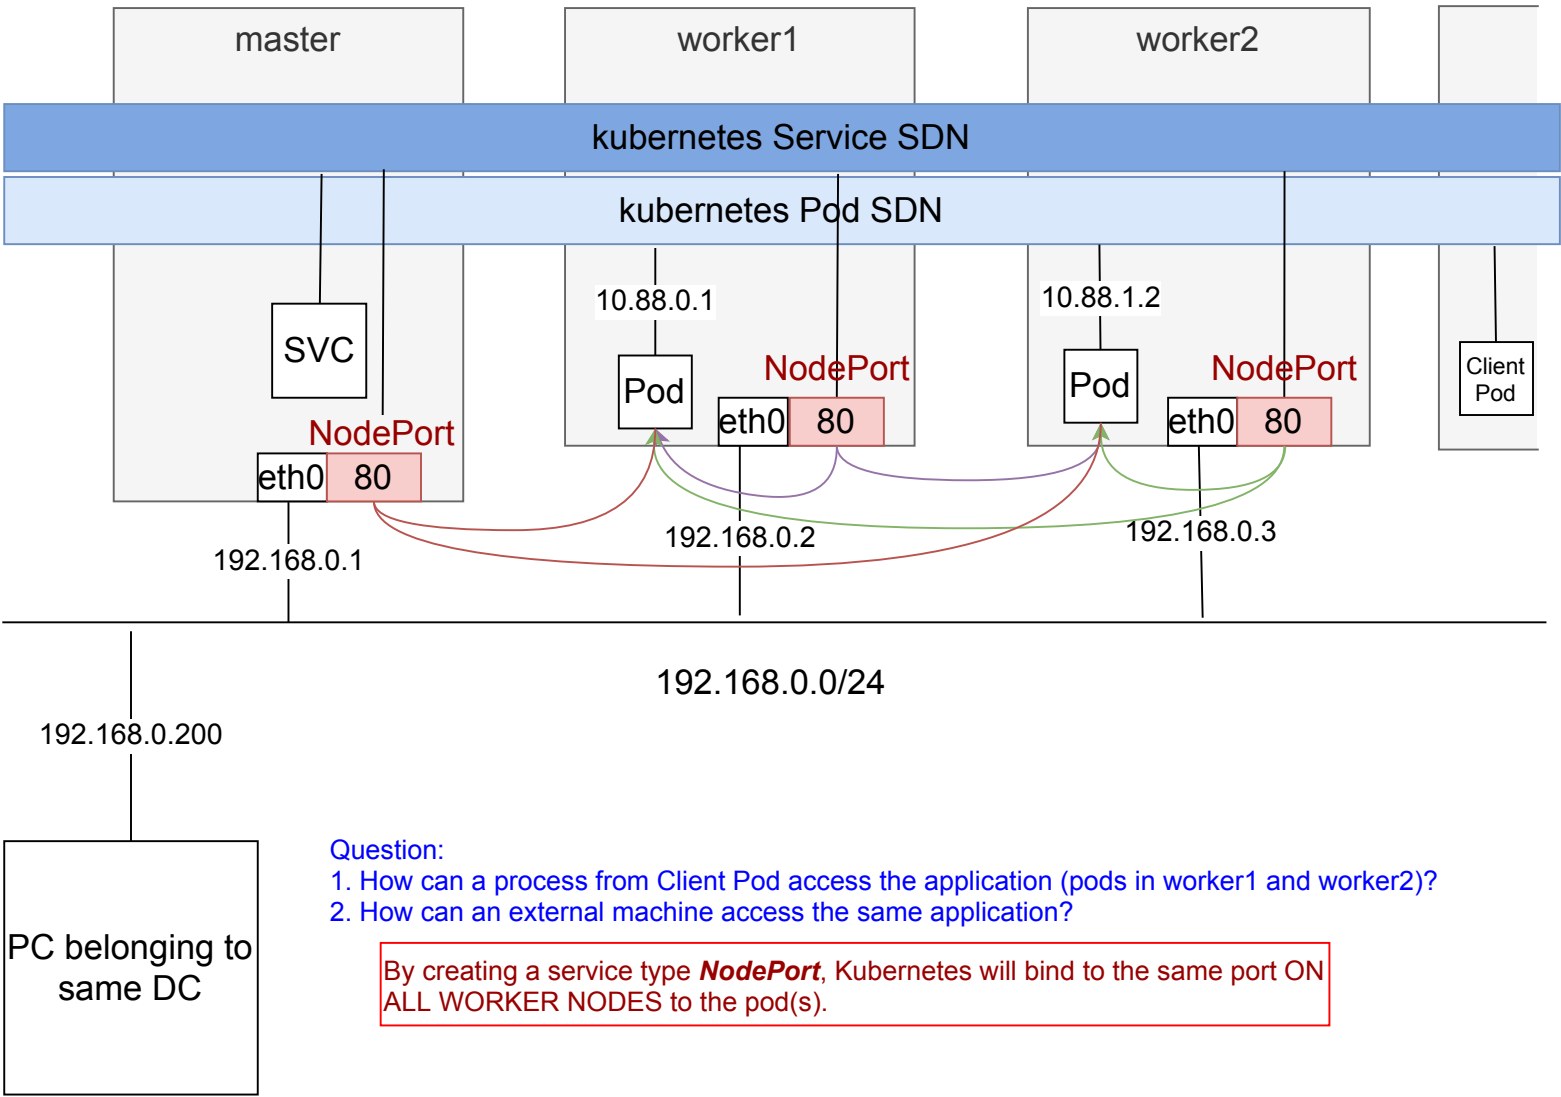
Warning: A service "can" refer to different pods, if the pods have the same label.

Sample of how Services are used



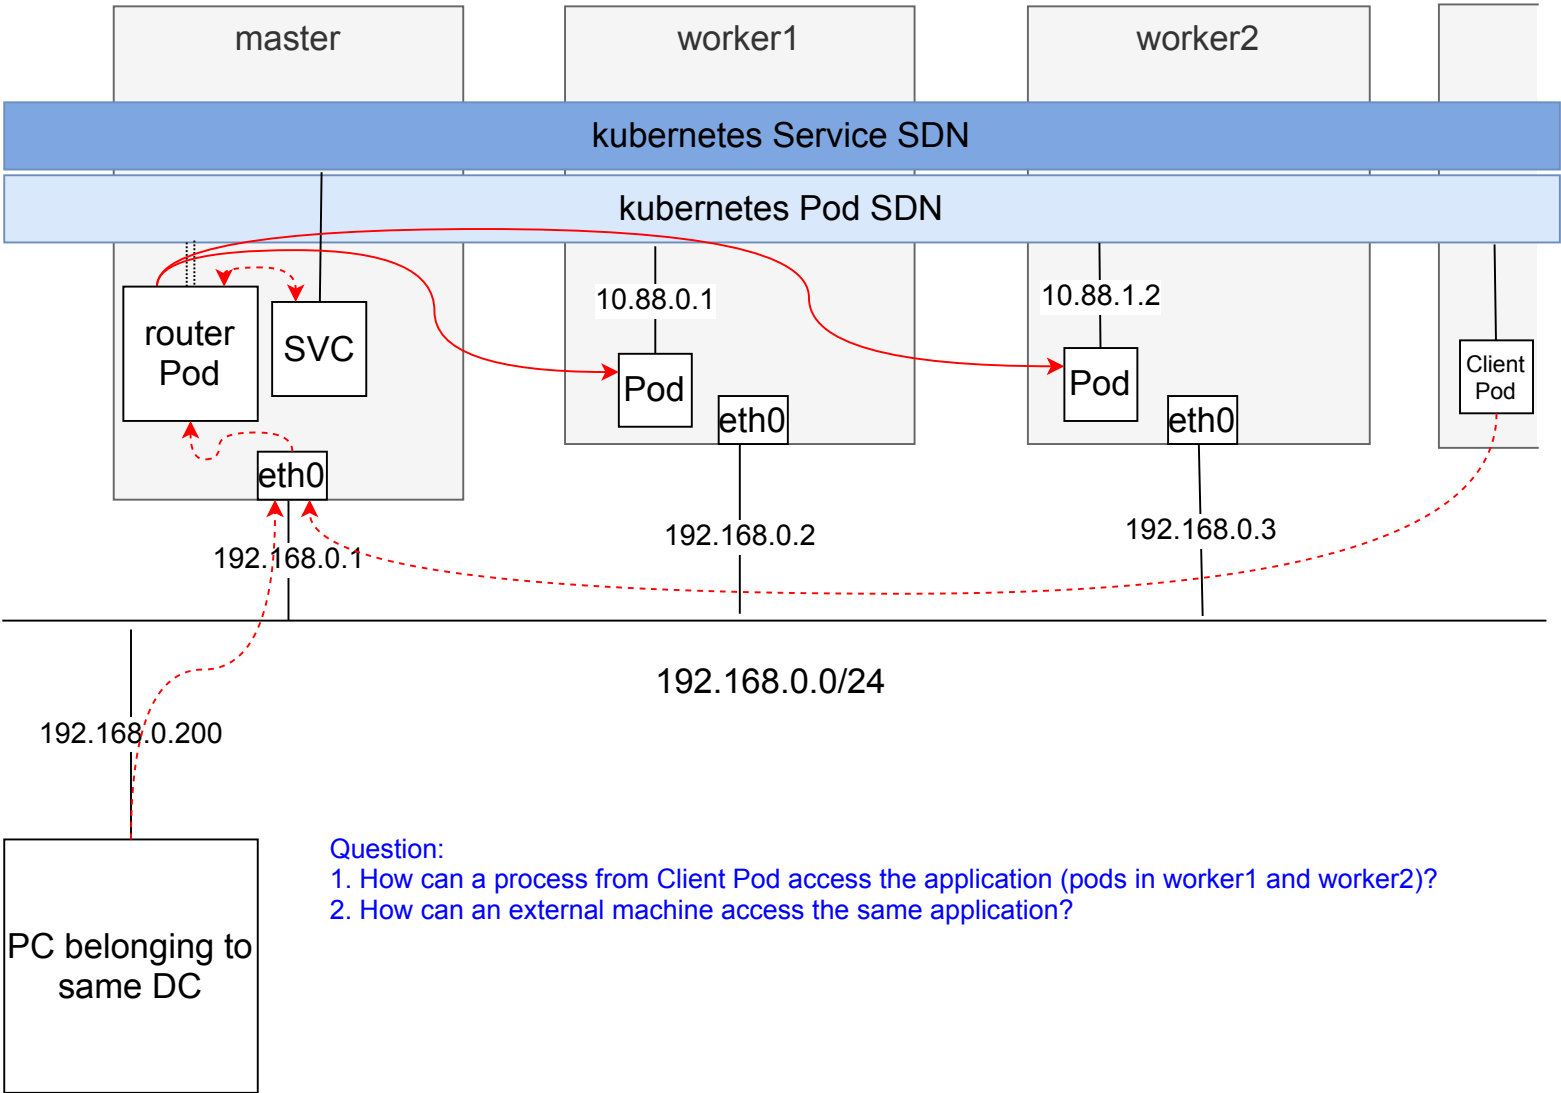


Question:
1. How can a process from Client Pod access the application (pods in worker1 and worker2)?
2. How can an external machine access the same application?



Question:
1. How can a process from Client Pod access the application (pods in worker1 and worker2)?
2. How can an external machine access the same application?

By creating a service type **NodePort**, Kubernetes will bind to the same port ON ALL WORKER NODES to the pod(s).



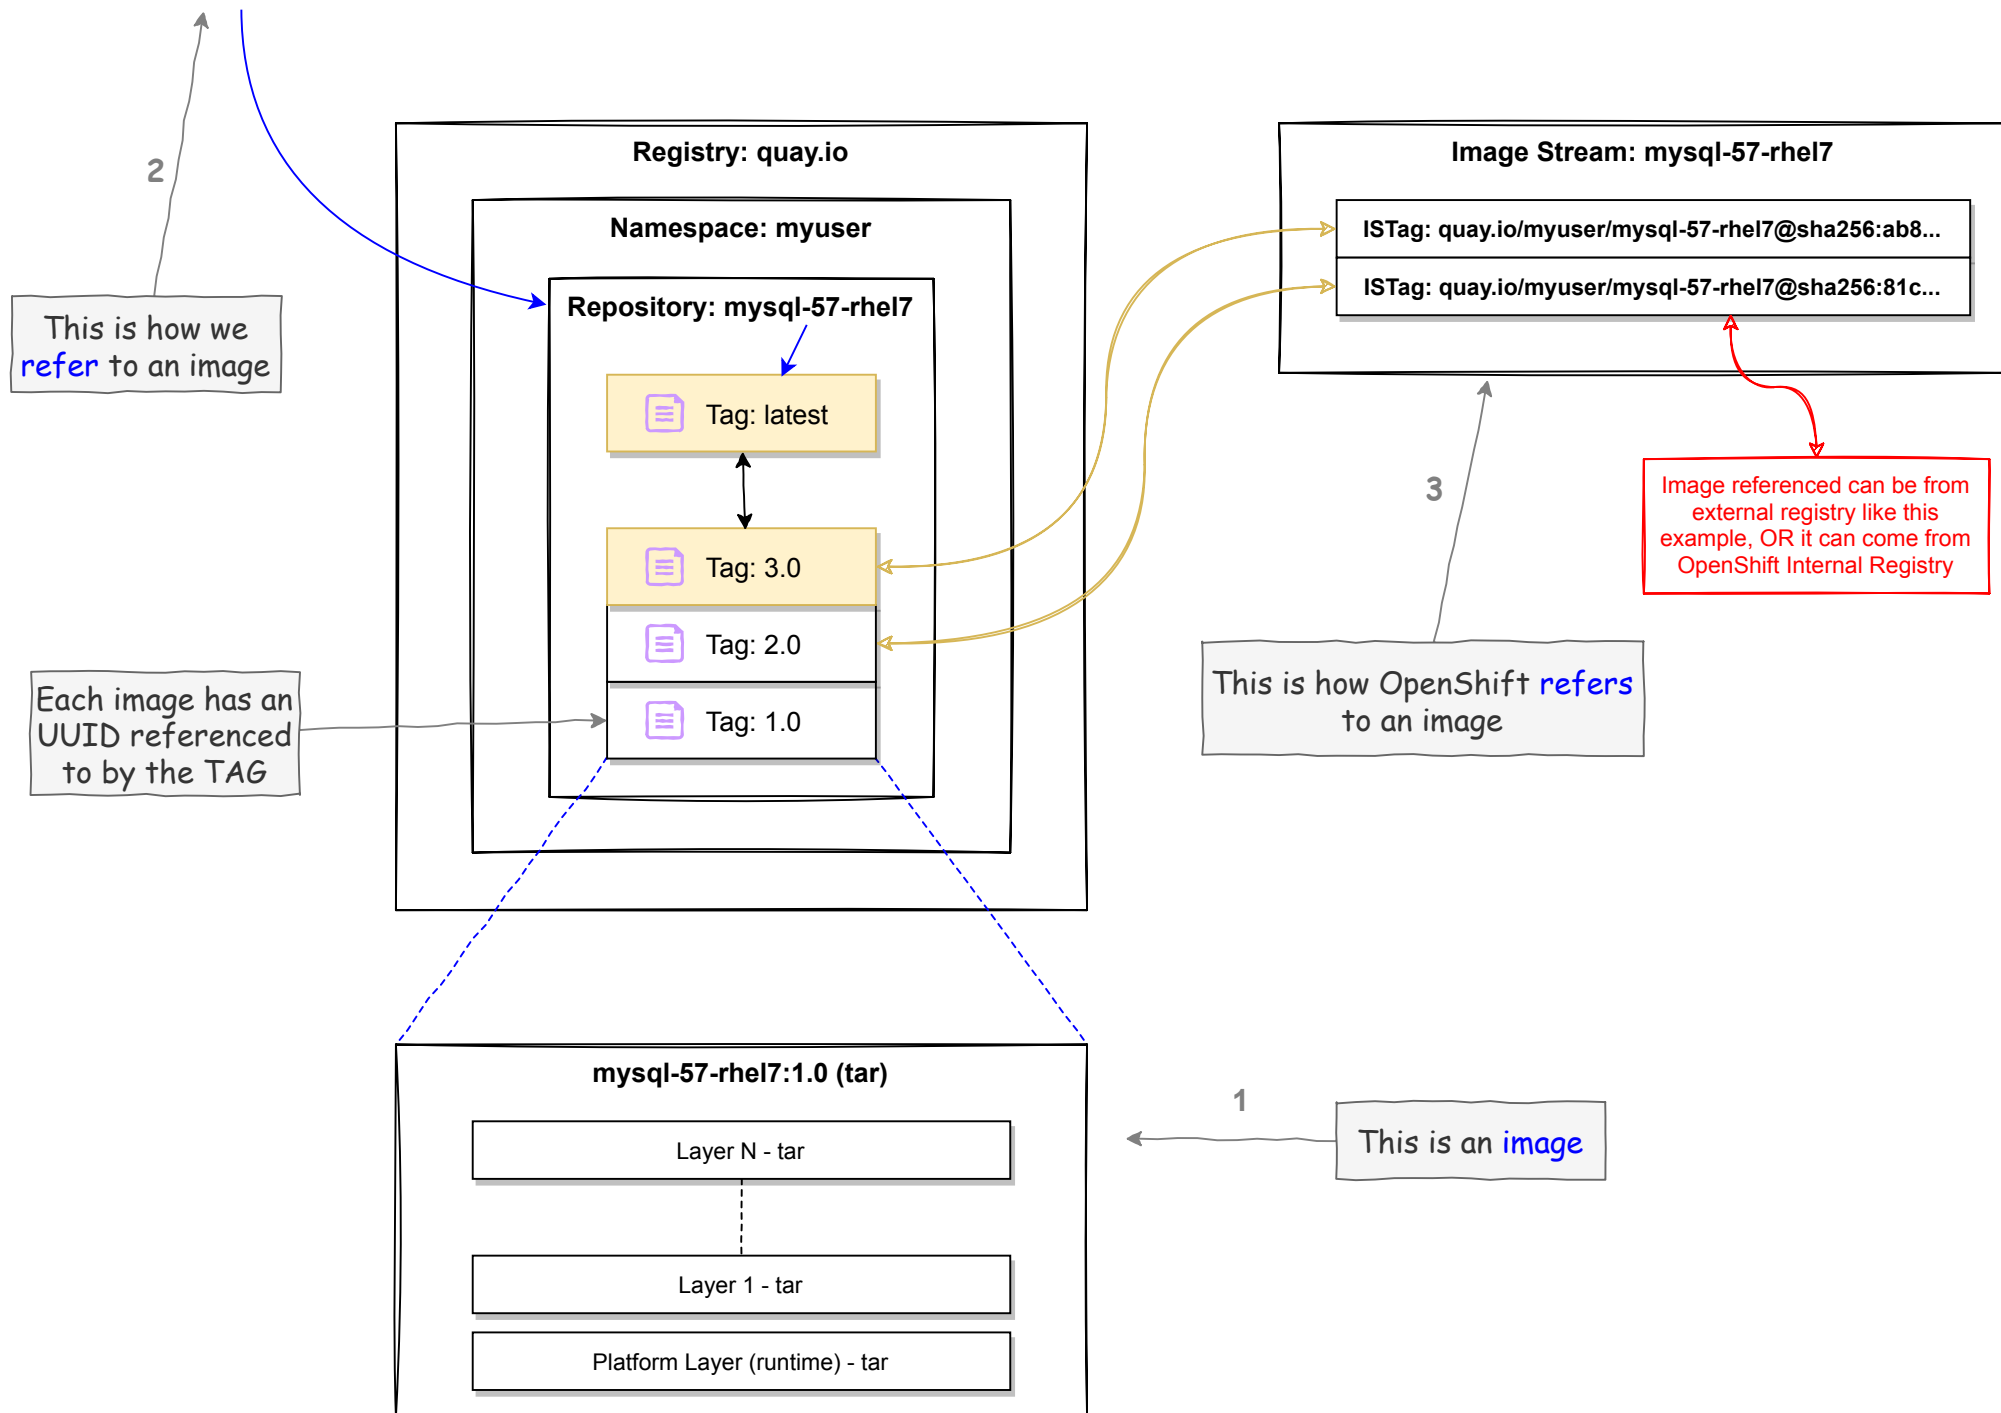
Question:
1. How can a process from Client Pod access the application (pods in worker1 and worker2)?
2. How can an external machine access the same application?

© 2020 Kelvin Lai



Container Image Naming Convention: <REGISTRY>[:<PORT>]/<NAMESPACE>/<REPOSITORY>[:<TAG>]

podman pull quay.io/myuser/mysql-57-rhel7



Deploying Applications with OpenShift

Methods to create applications:

1. Using existing containerised applications

```
oc new-app --docker-image=<IMAGE>
```

2. From Source Code using S2I

```
oc new-app <URL>
```

3. Using yaml/json file

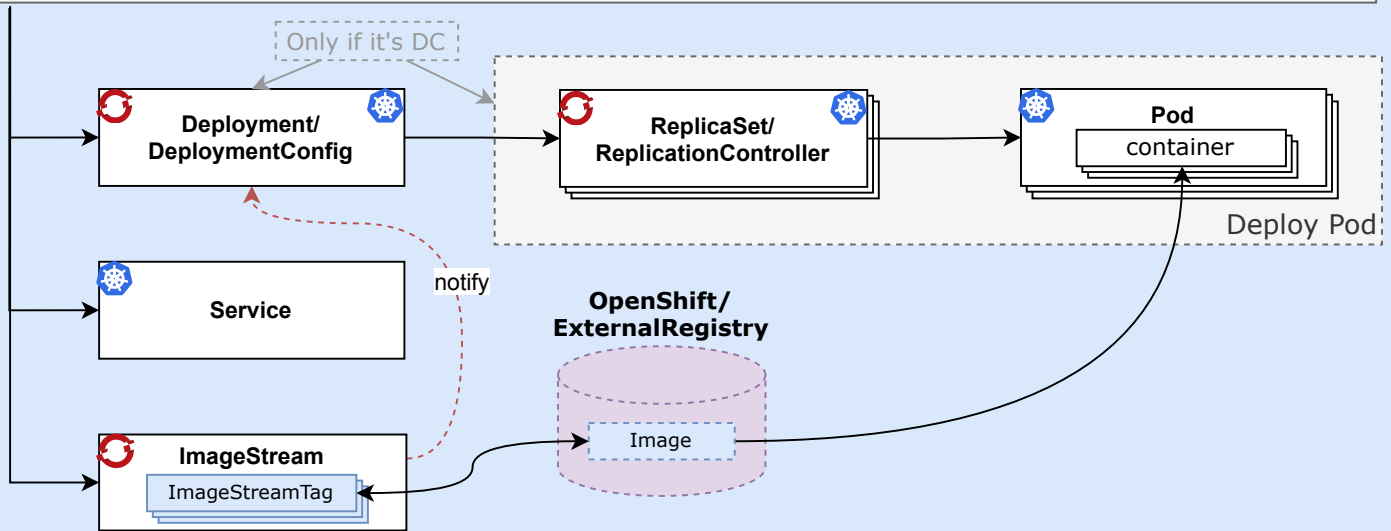
```
oc new-app -f <FILE>.yaml
```

4. Using template

```
oc new-app --template=<TEMPLATE> --param=<PARAM> --param-file=<PARAM_FILE>
```

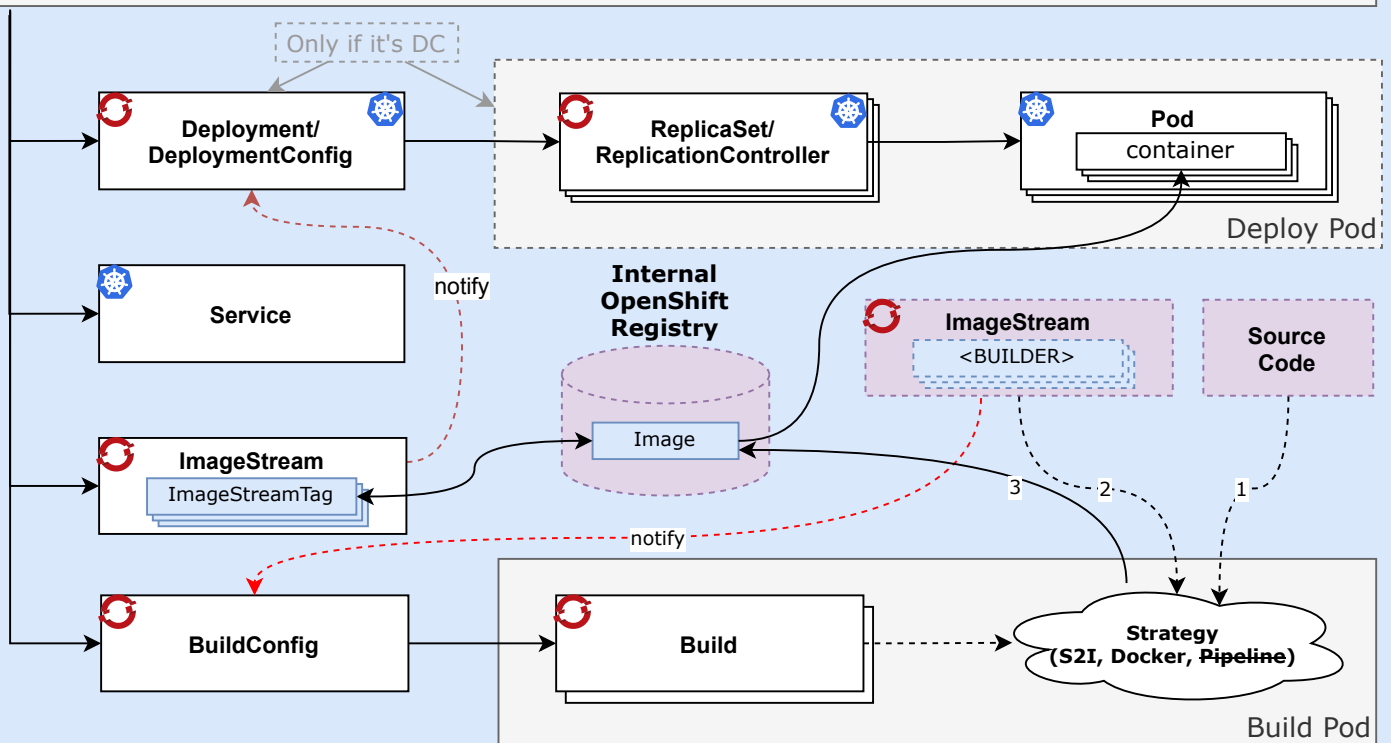
1. Use Existing Image

```
oc new-app [--as-deployment-config] [--docker-image <IMAGE>]
```



2. Managed Life Cycle

```
oc new-app [--as-deployment-config] [-i <BUILDER_IS>] [--strategy source|docker|pipeline] [--code <SRC_CODE>]
oc new-app [--as-deployment-config] [--strategy source|docker|pipeline] <BUILDER_IS> <SRC_CODE>
```



```
oc new-app -i php https://github.com/user/myapp#branch --context-dir <DIR>
```

```
oc new-app -i php:7.1 https://github.com/user/myapp
```

```
oc new-app php:7.1~https://github.com/user/myapp
```

NOTE: -i option needs git client to be installed

Options

-o json|yaml inspect resource definitions without creating

--name <NAME> adds a label "app=<NAME>" to all resources, Use `oc delete all -l "app=<NAME>"` to cleanup

IMPORT IMAGES

```
oc import-image <IMG_STREAM> [--confirm] --from <IMAGE> [--insecure]
where,
    <IMAGE> = <REGISTRY>[:<PORT>]/<NAMESPACE>/<REPOSITORY>[:<TAG>]
```

oc new-app command in OpenShift 4.5 makes use of *deployment* resource. Use --as-deployment-config if you wish to create *deployment config* instead.

SERVICE(SVC)

```
oc expose <DC/DEPLOYMENT/RC/RS/POD> <RESOURCE_NAME>
```

```
DNS NAME = <SVC>.<PROJ>[.svc.cluster.local]
ENVIRONMENT VARIABLE IN POD = <SVC>_SERVICE_HOST
```

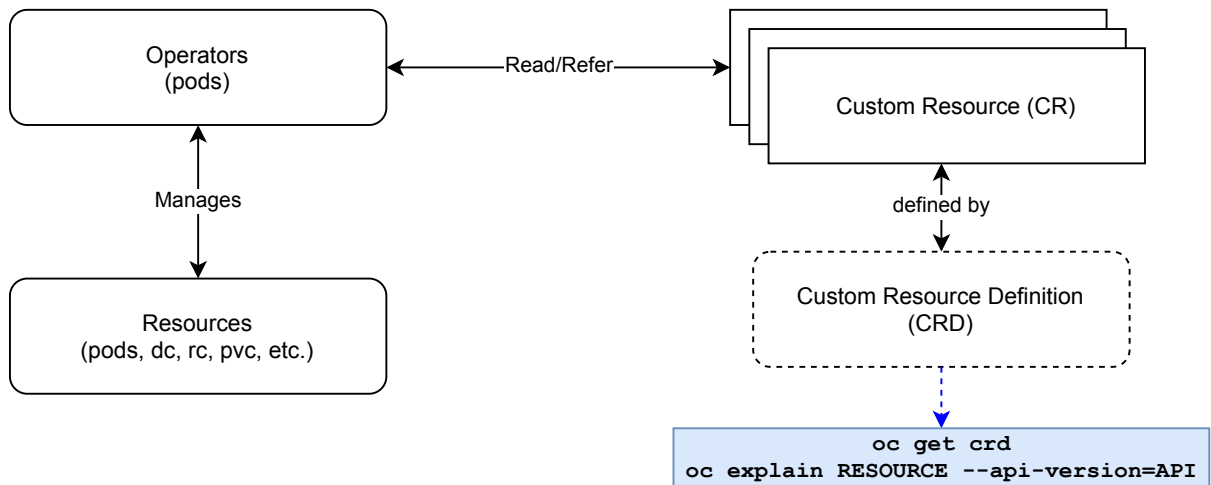


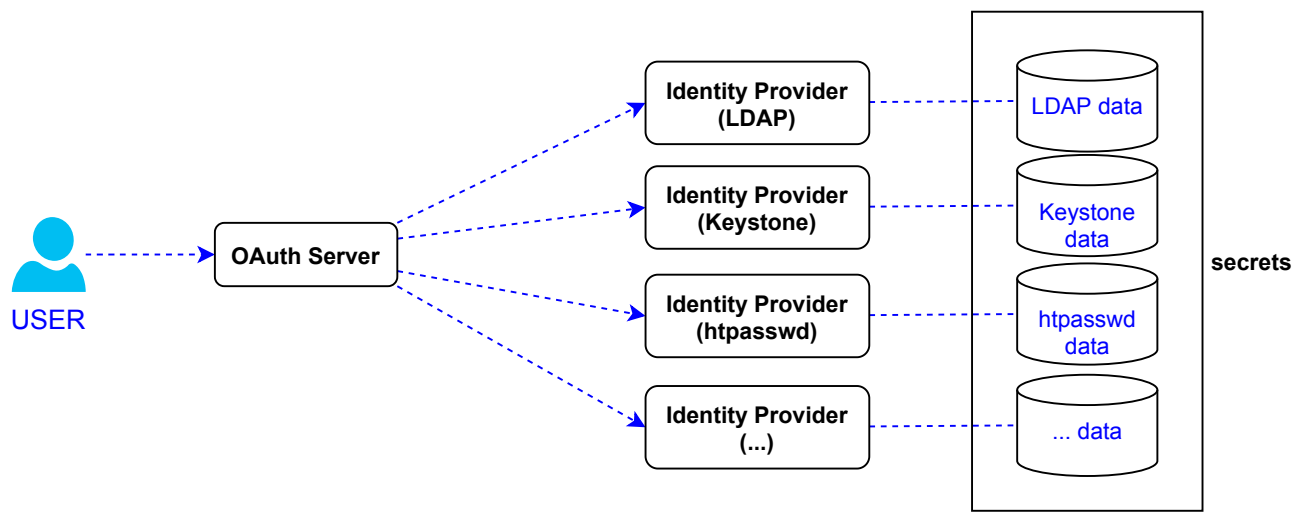
ROUTE

```
oc expose svc <SVC_NAME> [--name <ROUTE_NAME>] [--hostname <FQDN>]
```

```
DNS DEFAULT NAME = <ROUTE_NAME>-<PROJ>.<DOMAIN WILDCARD>
<DOMAIN_WILDCARD> = apps.<BASE_DOMAIN>
```



Operators






```
$ oc get co | grep auth
$ oc get pods -n openshift-authentication-operator
```

Project: openshift-authentication-operator




Manages/Configures



Project: openshift-authentication


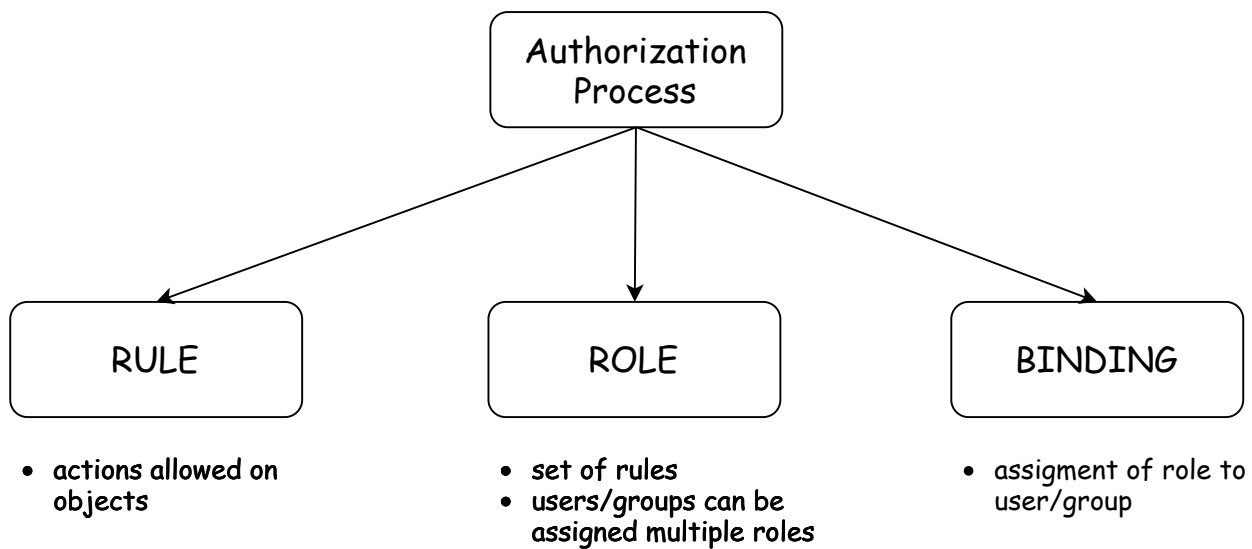
```
$ oc get pods -n openshift-authentication
```

```
$ oc get oauth cluster -o yaml
apiVersion: config.openshift.io/v1
kind: OAuth
metadata:
  name: cluster
  ...
spec:
  identityProviders:
  - htpasswd:
      fileName: htpasswd-secret
      mappingMethod: claim
      name: htpasswd_provider
      type: HTPasswd
```

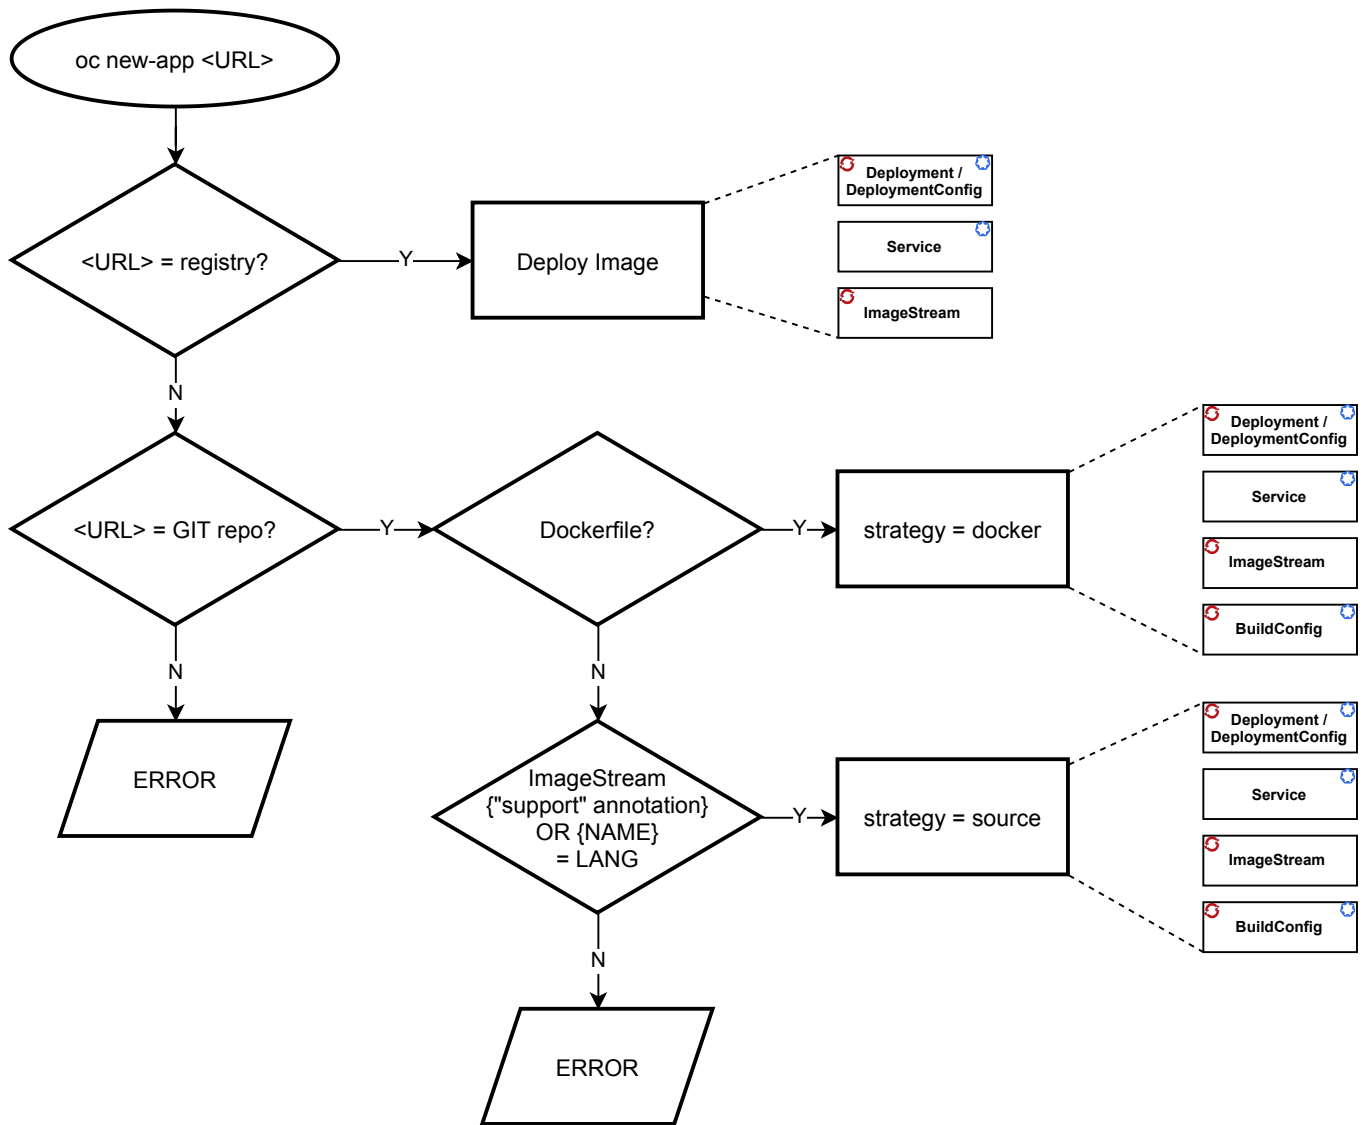


Project: openshift-config

```
$ htpasswd -cBb ./myusers <USER> <PASSWORD>
$ oc create secret generic htpasswd-secret --from-file htpasswd=./myhtpasswd -n openshift-config
```



oc new-app flowchart



Template

apiVersion: template.openshift.io/v1

kind: Template

metadata:

name: mytemplate

annotations:

description: "Description"

objects:

- apiVersion: v1

kind: Pod

metadata:

name: \${APP_NAME}

spec:

containers:

- env:

- name: ACCESS_CODE

value: \${APP_PASS}

image: superapp/hyperimage

name: myApp

ports:

- containerPort: 8080

protocol: TCP

parameters:

- description: Name of Pod

name: POD_NAME

value: myPod

required: true

- description: Application Secret Access Code

name: APP_PASS


generate: expression

from: "[a-zA-Z0-9]{8}"

labels:

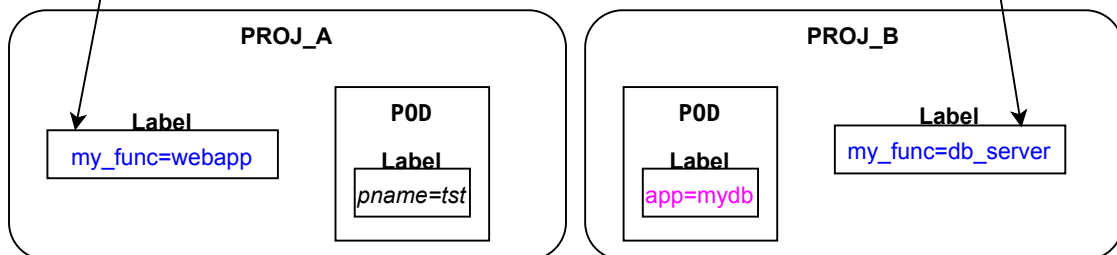
mylabel: myapp

Object Creation Order
↓



oc label namespace PROJ_A my_func=webapp

oc label namespace PROJ_B my_func=db_server

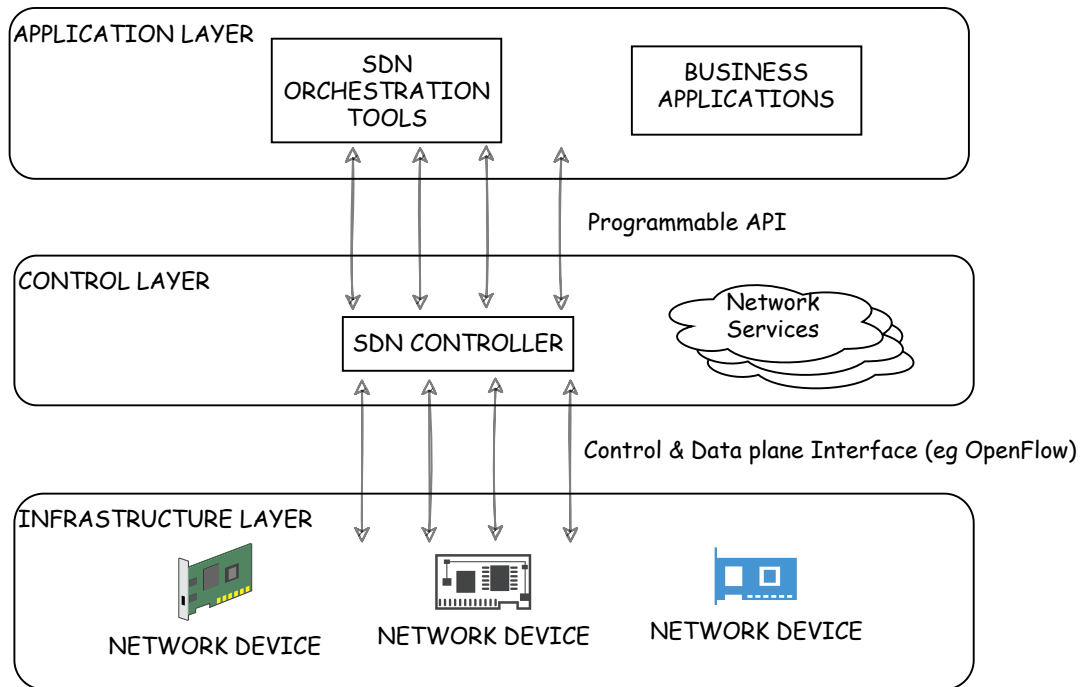


oc explain NetworkPolicy.spec

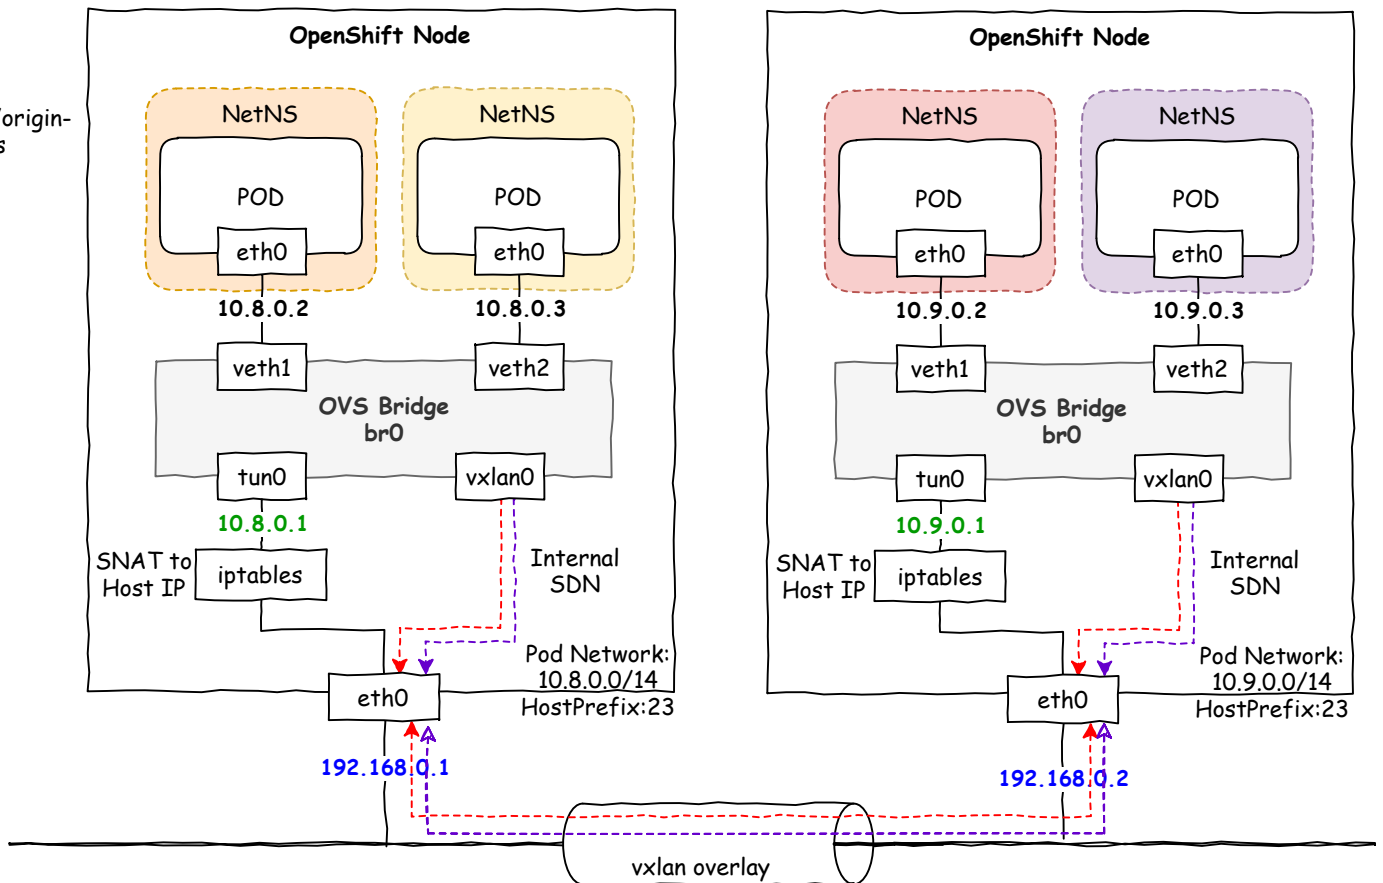
```
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: mynet_policy
spec:
  podSelector:
    matchLabels:
      app=mydb
  ingress:
    - from:
      - namespaceSelector:
          matchLabels:
            my_func: webapp
        podSelector:
          matchLabels:
            pname: tst
  ports:
    - port: 3306
      protocol: TCP
```

SDN

- Abstraction of network layers
- Decouple network control and forwarding functions




use
openshift/origin-
tools



nnnnnnnn . nnnnnnxx . xxxxxxxxx . xxxxxxxxx

POD Scheduling

1. Get a list of all NODES
2. Go through all the predicates for **FILTERing**. If NODE fails predicate rule, remove from list. Region affinity.
3. With remainder list of NODES, **prioritize** them using the weightage rules. NO filtering of NODES done here. Zone anti-affinity.
4. Select the NODE with highest points.



oc label node <NODE> <KEY>=<VALUE>
Region

<KEY> = failure-domain.beta.kubernetes.io/region

A set of hosts in closed geographical area. High speed connectivity.

Zone (availability zone)

<KEY> = failure-domain.beta.kubernetes.io/zone

A set of hosts that share common critical infra components (ups, switch, storage)

Upgrade Path Graph: https://access.redhat.com/labs/ocpupgradegraph/update_channel